

Overfitting and Underfitting

“No Free Lunch” Theorems

$Acc_G(L)$ = Generalization accuracy of learner L
= Accuracy of L on non-training examples
 \mathcal{F} = Set of all possible concepts, $y = f(\mathbf{x})$

Theorem: For any learner L , $\frac{1}{|\mathcal{F}|} \sum_{\mathcal{F}} Acc_G(L) = \frac{1}{2}$
(given any distribution \mathcal{D} over \mathbf{x} and training set size n)

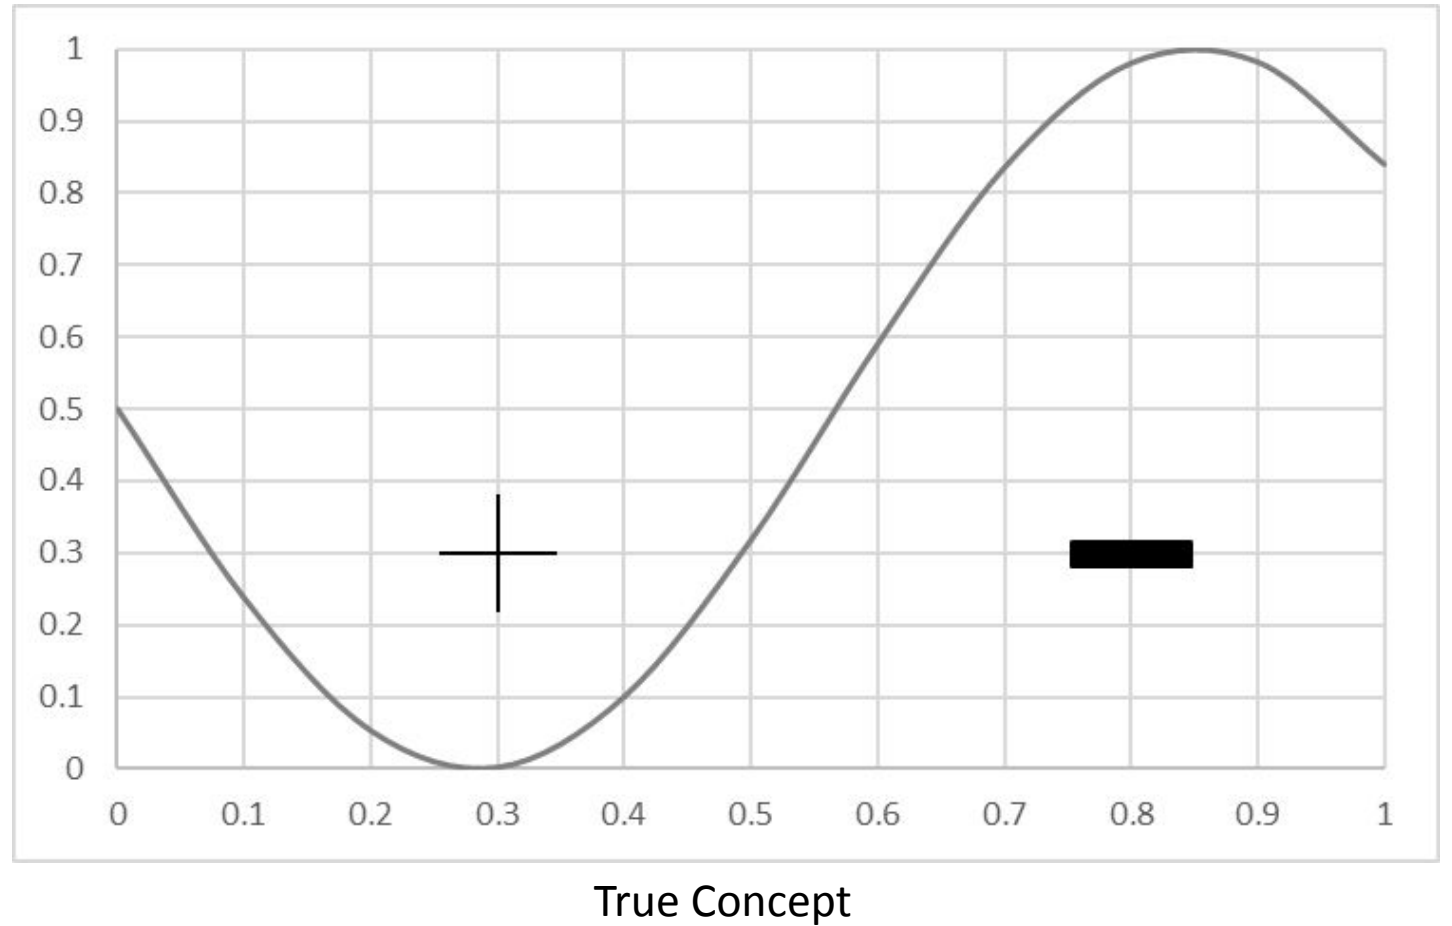
Bias and Variance

- Bias – error caused because the model can not represent the concept
- Variance – error caused because the learning algorithm overreacts to small changes (noise) in the training data

$$\text{TotalLoss} = \text{Bias} + \text{Variance} (+ \text{noise})$$


Visualizing Bias

- Goal: produce a model that matches this concept



Visualizing Bias

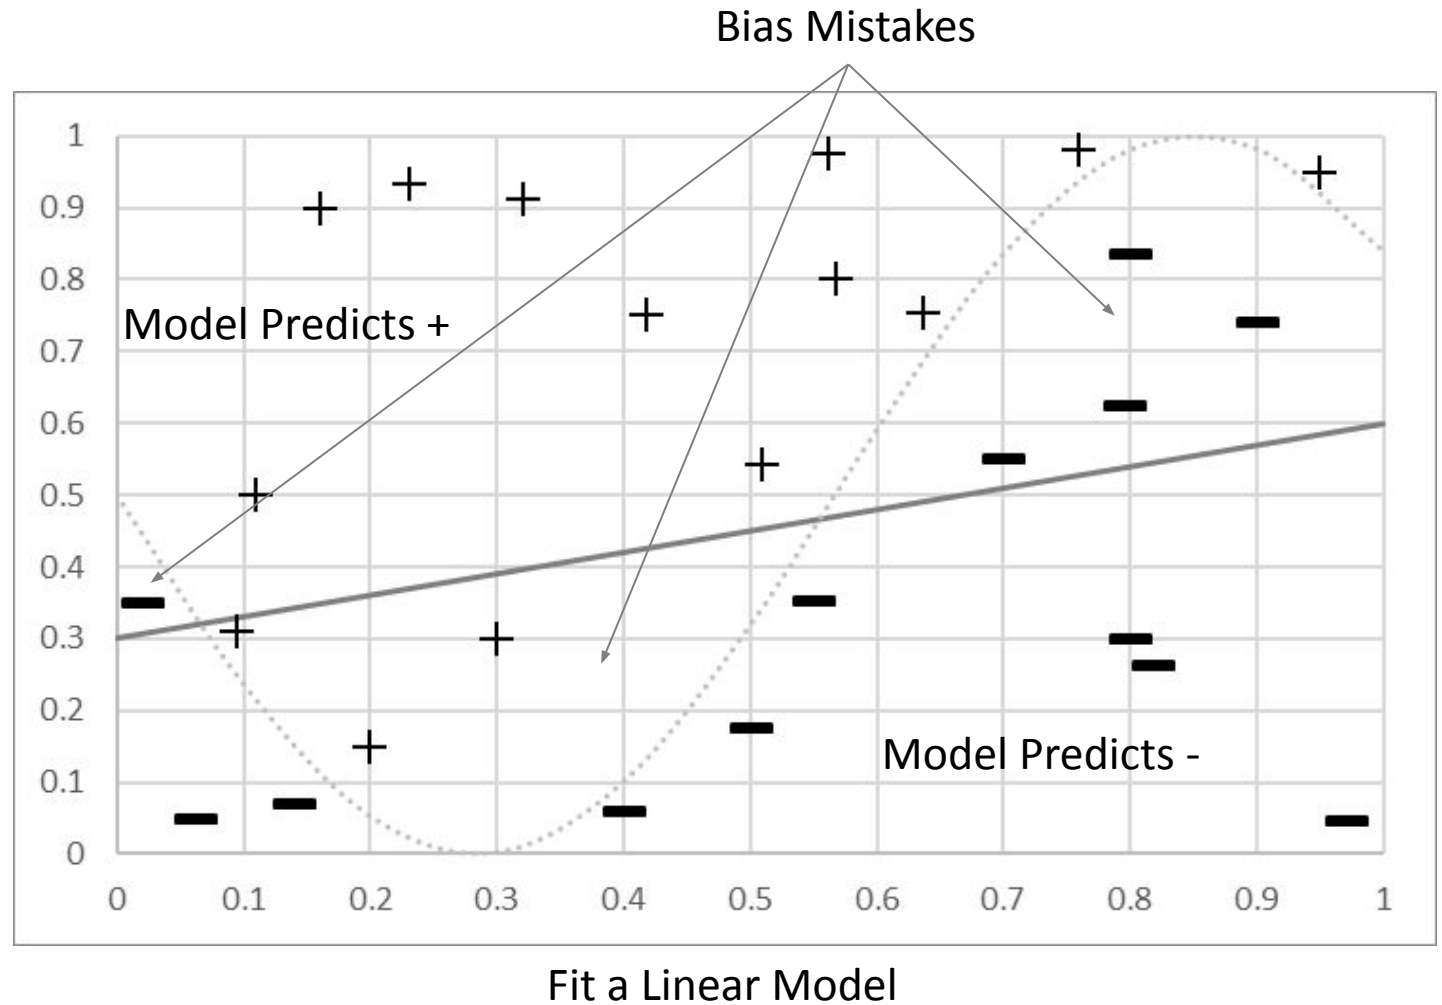
- Goal: produce a model that matches this concept
- Training Data for the concept



Training Data


Visualizing Bias

- Goal: produce a model that matches this concept
- Training Data for concept
- Bias: Can't represent it...




Visualizing Bias

- Goal: produce a model that matches this concept
- New data, new model




Visualizing Variance

- Goal: produce a model that matches this concept
- New data, new model
- New data, new model...
- Variance: Sensitivity to changes & noise

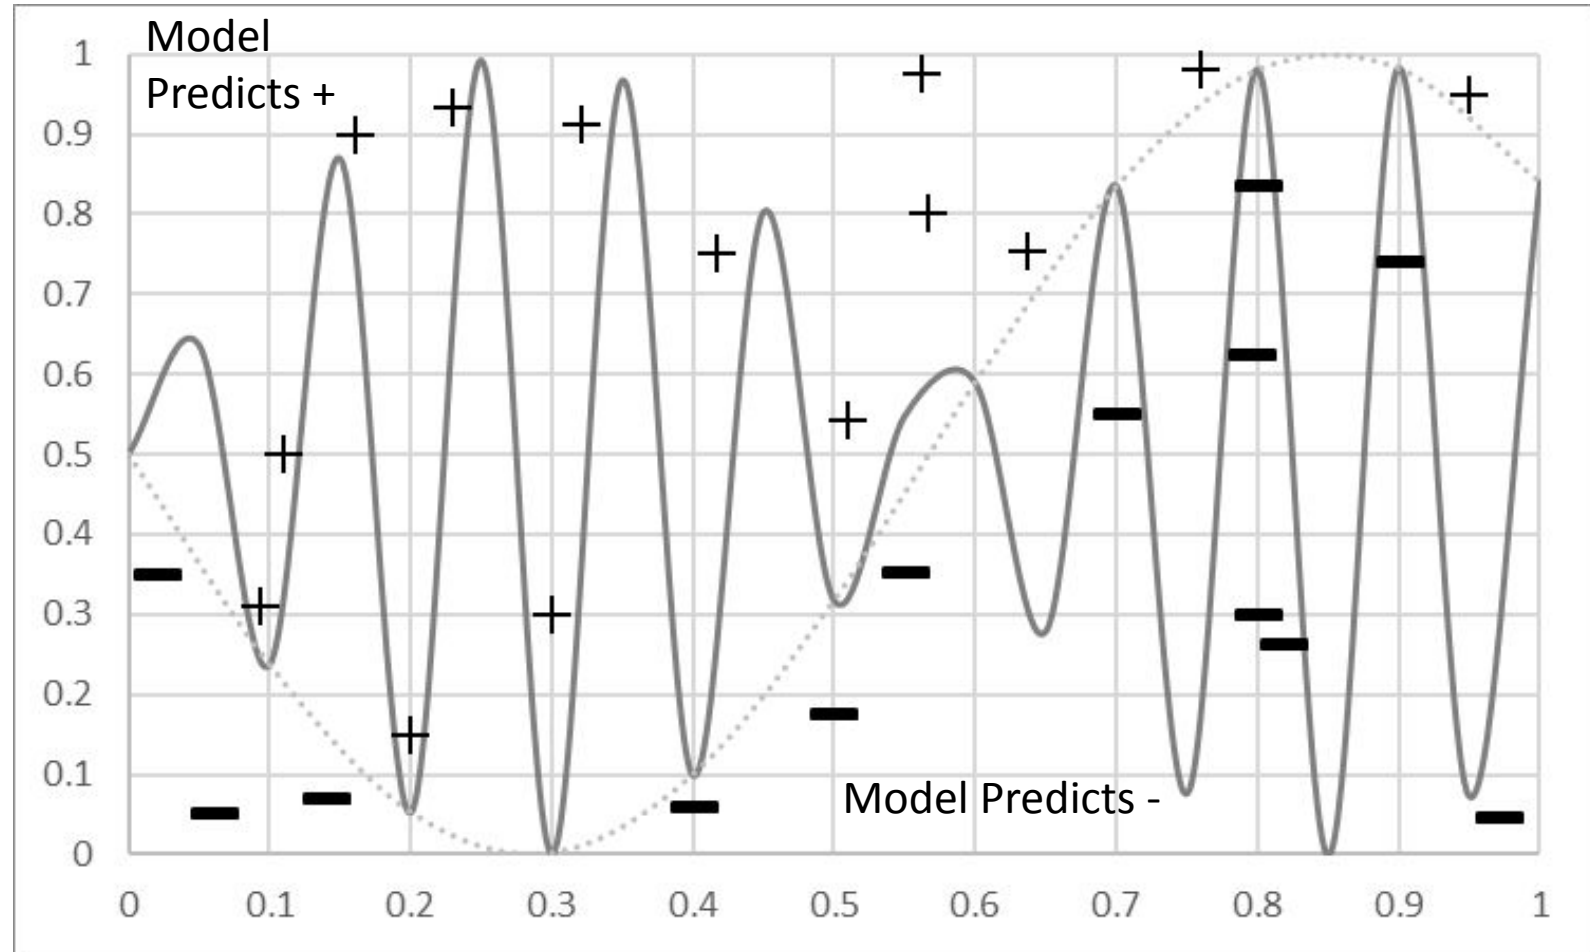


Another way to think about Bias & Variance




Bias and Variance: More Powerful Model

- Powerful Models can represent complex concepts
- No Mistakes!



Bias and Variance: More Powerful Model


- But get more data...
- Not good!



Overfitting vs Underfitting


Overfitting

- Fitting the data too well
 - Features are noisy / uncorrelated to concept
 - Modeling process very sensitive (powerful)
 - Too much search




Underfitting



- Learning too little of the true concept
 - Features don't capture concept
 - Too much bias in model
 - Too little search to fit model




The Effect of Noise



The Effect of Features



Throw out x_2



New x_3

- Not much info
- Won't learn well
- Powerful -> high variance

- Captures concept
- Simple model -> low bias
- Powerful -> low variance

The Power of a Model Building Process


Weaker Modeling Process (higher bias)

- Simple Model (e.g. linear)
- Fixed sized Model (e.g. fixed # weights)
- Small Feature Set (e.g. top 10 tokens)
- Constrained Search (e.g. few iterations of gradient descent)

More Powerful Modeling Process (higher variance)

- Complex Model (e.g. high order polynomial)
- Scalable Model (e.g. decision tree)
- Large Feature Set (e.g. every token in data)
- Unconstrained Search (e.g. exhaustive search)

Example of Under/Over-fitting



Ways to Control Decision Tree Learning

- Increase minToSplit
- Increase minGainToSplit
- Limit total number of Nodes
- Penalize complexity

$$Loss(S) = \sum_i^n Loss(y_i^{\wedge}, y_i) + \alpha Log_2(\# Nodes)$$

Ways to Control Logistic Regression

- Adjust Step Size
- Adjust Iterations / stopping criteria of Gradient Descent
- Regularization

$$Loss(S) = \sum_i^n Loss(y_i^{\wedge}, y_i) + \alpha \overset{\# Weights}{\sum_j} |w_j|$$

Modeling to Balance Under & Overfitting

- Data
 - Learning Algorithms
 - Feature Sets
 - Complexity of Concept
 - Search and Computation
-
- Parameter sweeps!

Parameter Sweep

optimize first parameter

for p in [setting_certain_to_underfit, ..., setting_certain_to_overfit]:

do cross validation to estimate accuracy

find the setting that balances overfitting & underfitting

optimize second parameter

etc...

examine the parameters that seem best and adjust whatever you can...

Types of Parameter Sweeps

- Optimize one parameter at a time
 - Optimize one, update, move on
 - Iterate a few times
- Gradient descent on meta-parameters
 - Start somewhere 'reasonable'
 - Computationally calculate gradient wrt change in parameters
- Grid
 - Try every combination of every parameter
- Quick vs Full runs
 - Expensive parameter sweep on 'small' sample of data (e.g. grid)
 - A bit of iteration on full data to refine
- Intuition & Experience
 - Learn your tools
 - Learn your problem

Summary of Overfitting and Underfitting

- Bias / Variance tradeoff a primary challenge in machine learning
- Internalize: More powerful modeling is not always better
- Learn to identify overfitting and underfitting
- Tuning parameters & interpreting output correctly is key